Rohan Taori

http://rohantaori.com/ | rohantaori@berkeley.edu | San Francisco Bay Area

Education

University of California, Berkeley   Aug '16 - May '19
  B.S. in Electrical Engineering and Computer Science   GPA: 3.98
  Awards: Regents' & Chancellor's Scholar (top 2% of students), Highest Honors & Dean's List (top 3% of engineers)
  Memberships: Eta Kappa Nu (HKN) honor society, 2019 Accel Scholar (VC firm Accel and Berkeley)

Research & Work Experience

Uber ATG – AI Research Resident   Sep '19 - *Present*
  • Working with Prof. Raquel Urtasun on improving the reliability of autonomous driving systems

Berkeley AI Research – Undergraduate Researcher   Aug '16 - *Present*
  • Worked with Prof. Ben Recht on understanding robustness in real-world distribution shifts
    Study was the largest to date: compared 150+ models with 200+ ImageNet distribution shifts
  • Worked with Prof. Jitendra Malik on sample efficient deep RL using non-parametric value functions
  • Worked with Prof. Carlo Séquin on finding optimal solutions of irregular solids via gradient descent

NVIDIA – Deep Learning Research Intern   May - Aug '19
  • Worked with Bryan Catanzaro on new techniques for large-scale texture synthesis in near real-time
    Developed a new architecture which improves generalization to unseen test-time textures

Facebook – Software Engineering Intern   May - Aug '18
  • Built infrastructure to allow bots to send messages to hundreds of millions of users on the platform

Salesforce – Software Engineering Intern   May - Aug '17
  • Developed an API tool to allow external developers to integrate with the Salesforce database via REST calls

Teaching

Machine Learning @ Berkeley – VP of Education   May '17 - Dec '18
  • Led a 15-person team running two student classes, workshops, member bootcamps, and weekly reading groups
  • Created material for new deep learning class of 200 students; spearheaded creation of machine learning workshop series
  • Goal to bring ML education to students: content for fall17, spring18, fall18, member class, workshops

Papers

Rohan Taori, Achal Dave, Vaishaal Shankar, Nicholas Carlini, Benjamin Recht, Ludwig Schmidt. *When Robustness Doesn't Promote Robustness: Synthetic vs. Natural Distribution Shifts on ImageNet.* In submission to ICLR 2020. link.

Guilin Liu, **Rohan Taori**, Zhiding Yu, Ting-Chung Wang, Edward Liu, Karan Sapra, Fitsum Reda, Brandon Rowlett, Andrew Tao, Bryan Catanzaro. *Real-Time Universal Texture Synthesis using Feature Maps as Transpose Convolution Filters*. In submission to CVPR 2020. link.

Rohan Taori, Amog Kamsetty, Brenton Chu, Nikita Vemuri. *Targeted Adversarial Examples for Black Box Audio Systems.* Published at the 2019 IEEE Security and Privacy Workshops, pp. 15-20. Appeared in NeurIPS 2018 Security in ML Workshop and DEFCON 26 CAAD Village. link.

Murtaza Dalal*, Alexander Li*, **Rohan Taori***. *Autoregressive Models: What Are They Good For?* Appearing in NeurIPS 2019 Information Theory and ML Workshop. link.

Carlo Séquin, Yifat Amir, Ruta Jawale, Hong Jeon, Alex Romano, **Rohan Taori**. *Modular Toroids Constructed from Nonahedra*. Berkeley EECS 2017 Tech Report. link.

Project Highlights

Take a Picasso – CalHacks 3.0 Project   Oct '16
  • Robotic sketch artist; won Best Hardware Hack, Best 3D Printed Hack, and grant from Peter Thiel's 1517 Fund

MyFi.py – Berkeley EECS Hackathon   Sep '16
  • Built a network monitoring platform to inform users about traffic patterns and local events; won 2nd place