Rohan Taori

http://rohantaori.com/ | rohantaori@berkeley.edu | San Francisco Bay Area

Education

University of California, Berkeley Aug '16 - May '19

B.S. in Electrical Engineering and Computer Science GPA: 3.98

Awards: Regents' & Chancellor's Scholar (top 2% of students), Highest Honors & Dean's List (top 3% of engineers)

Memberships: Eta Kappa Nu (HKN) honor society, 2019 Accel Scholar (VC firm Accel and Berkeley)

Classes: Machine Learning, Artificial Intelligence, Neural Networks, Deep Unsupervised Learning*, Real Analysis, Probability & Random Processes, Algorithms, Operating Systems, Computer Security, Data Structures

Research & Work Experience

Uber ATG – AI Research Resident Sep '19 - *Present*

- Working with Prof. Raquel Urtasun on improving the reliability of autonomous driving systems.

Berkeley AI Research – Undergraduate Researcher Aug '16 - *Present*

- Worked with Prof. Ben Recht on understanding robustness in real-world distribution shifts.
  - Study was the largest to date: compared 150+ models with 200+ ImageNet distribution shifts.
- Worked with Prof. Jitendra Malik on sample efficient deep RL using non-parametric value functions.
- Worked with Prof. Carlo Séquin on finding optimal solutions of irregular solids via gradient descent.

NVIDIA – Deep Learning Research Intern May - Aug '19

- Worked with Dr. Bryan Catanzaro on new techniques for texture synthesis.

Facebook – Software Engineering Intern May - Aug '18

- Built infrastructure to allow bots to send messages to hundreds of millions of users on the platform.

Salesforce – Software Engineering Intern May - Aug '17

- Developed an API tool to allow external developers to integrate with the Salesforce database via REST calls.

Teaching

Machine Learning @ Berkeley – VP of Education May '17 - Dec '18

- Led a 15-person team running two student classes, workshops, member bootcamps, and weekly reading groups.
- Created material for new deep learning class of 200 students; spearheaded creation of machine learning workshop series.
- Goal to bring ML education to students: content for fall17, spring18, fall18, member class, workshops.

Papers

Rohan Taori, Achal Dave, Vaishaal Shankar, Nicholas Carlini, Benjamin Recht, Ludwig Schmidt. ***When Robustness Doesn't Promote Robustness: Synthetic vs. Natural Distribution Shifts on ImageNet.*** In submission to ICLR 2020.

Rohan Taori, Amog Kamsetty, Brenton Chu, Nikita Vemuri. ***Targeted Adversarial Examples for Black Box Audio Systems.*** Published at the 2019 IEEE Security and Privacy Workshops, pp. 15-20. Appeared in NeurIPS 2018 Security in ML Workshop and DEFCON 26 CAAD Village. link.

Murtaza Dalal*, Alexander Li*, **Rohan Taori***. ***Autoregressive Models: What Are They Good For?*** Appearing in NeurIPS 2019 Information Theory and ML Workshop. link.

Carlo Séquin, Yifat Amir, Ruta Jawale, Hong Jeon, Alex Romano, **Rohan Taori**. ***Modular Toroids Constructed from Nonahedra***. Berkeley EECS 2017 Tech Report. link.

Work from NVIDIA internship (second-author paper) in submission to CVPR 2020.

Project Highlights

Take a Picasso – CalHacks 3.0 Project Oct '16

- Robotic sketch artist; won Best Hardware Hack, Best 3D Printed Hack, and grant from Peter Thiel's 1517 Fund

MyFi.py – Berkeley EECS Hackathon Sep '16

- Built a network monitoring platform to inform users about traffic patterns and local events; won 2nd place